


Introduction to Kubernetes

[Waheed Iqbal](#)

Infrastructure Management for Scalable Applications (Fall 2022)
Department of Data Science, FCIT, University of the Punjab
Lahore, Pakistan






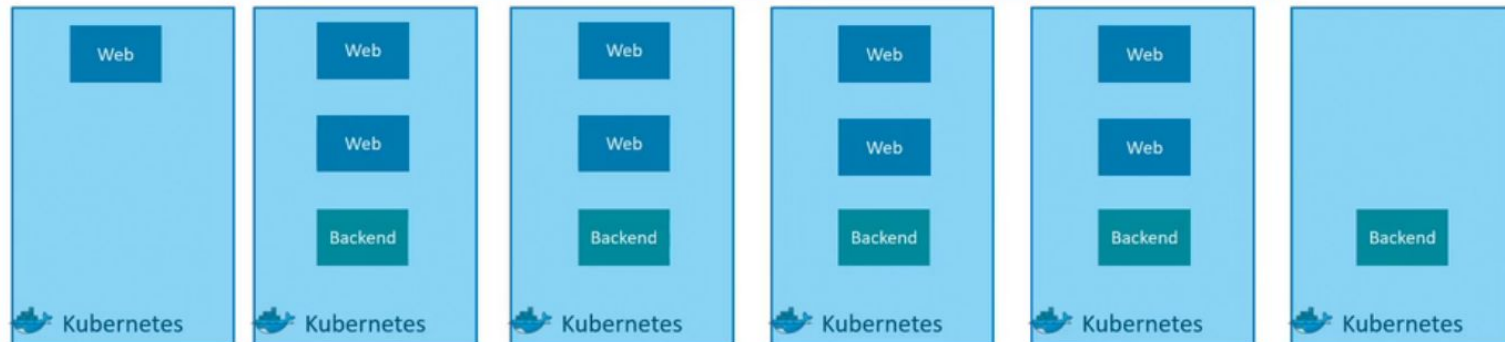
Container Orchestration

- Container orchestration is the **automated management** of containerized applications
- Orchestration platforms automate the **deployment, scaling, and management of containerized applications across multiple hosts**
- They offer features like **service discovery, load balancing, storage orchestration, and automated rollouts and rollbacks**
- Popular container orchestration platforms include **Kubernetes, Docker Swarm, and Apache Mesos**
- Container orchestration is a key component of modern DevOps and **cloud-native application development**

Orchestration






Orchestration



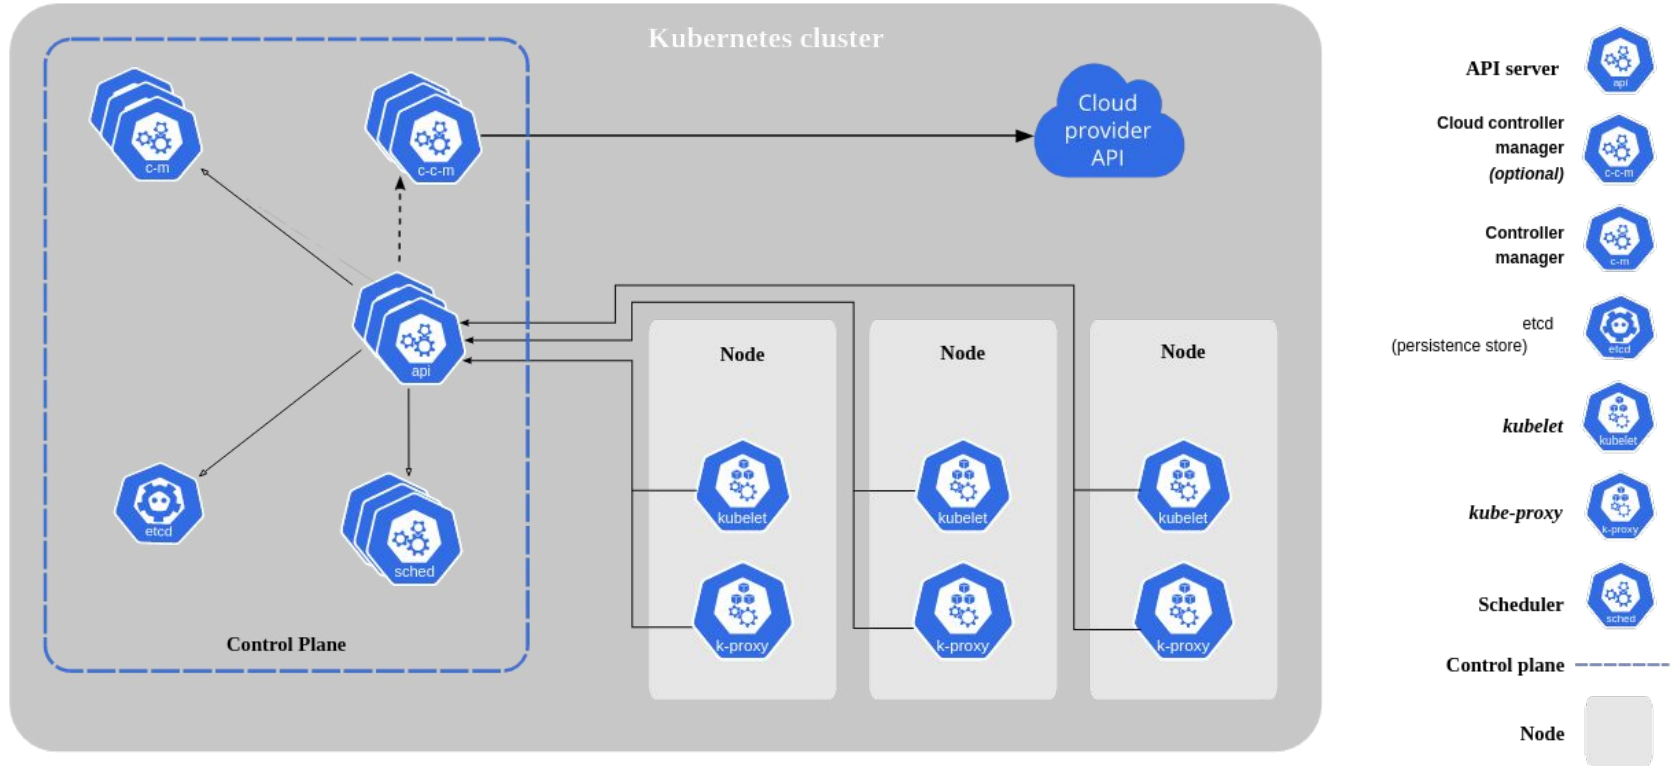
Kubernetes

- Kubernetes is an open-source **container orchestration platform** that automates the deployment, scaling, and management of containerized applications
- It was originally developed by **Google** and is now maintained by the **Cloud Native Computing Foundation** (CNCF)
- Kubernetes is based on a **distributed architecture** and uses a **master-worker node model** to manage and orchestrate containers across multiple hosts
- Kubernetes can be deployed on any cloud platform or **on-premises infrastructure**

Kubernetes Components




Kubernetes Components (Cont.)



Pod

- Pod is a collection of **one or more Linux containers**, packaged together to maximize the benefits of resource sharing via cluster management
- Pods provide a **logical host for containers** and are used to encapsulate and run containerized applications and services
- Containers within a Pod **share the same network namespace**, which means they can communicate with each other using localhost and share the same IP address and port space
- Containers within a Pod also share the **same storage namespace**, which means they can access and use the same storage volumes mounted into the Pod
- Pods can be **scaled horizontally** by creating multiple replicas of the same Pod

Pod (Cont.)



Pod (Cont.)

```
kubectl run nginx --image nginx
```


```
kubectl get pods
```

```
C:\Kubernetes>kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-8586cf59-whssr	0/1	ContainerCreating	0	3s

```
C:\Kubernetes>kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
nginx-8586cf59-whssr	1/1	Running	0	8s



Minikube

- Minikube is a lightweight, **single-node Kubernetes cluster** on a local machine, which enables developers to test and experiment with Kubernetes without having to set up a full-scale production environment.
- Minikube is **easy to install** and can be run on most operating systems, including Windows, macOS, and Linux.
- Minikube uses a single-node cluster configuration, which means that it runs all the Kubernetes components, such as the API server, etcd, and kubelet, on a single virtual machine or container, instead of across multiple nodes.
- Minikube provides a **local Docker registry**, which enables developers to build and test container images on the same machine where the Kubernetes cluster is running.
- Minikube can be used to test and debug Kubernetes applications and services in a local environment, before deploying them to a production cluster.

Installing Minikube

- `curl -LO https://storage.googleapis.com/minikube/releases/latest/minikube_latest_amd64.deb`
- `sudo dpkg -i minikube_latest_amd64.deb`
- If everything goes well, execute minikube start command.

```
~$ minikube start
🐹 minikube v1.29.0 on Ubuntu 22.04
🌟 Using the virtualbox driver based on existing profile
👉 Starting control plane node minikube in cluster minikube
🔄 Restarting existing virtualbox VM for "minikube" ...
📡 Preparing Kubernetes v1.26.1 on Docker 20.10.23 ...
🔗 Configuring bridge CNI (Container Networking Interface) ...
🔍 Verifying Kubernetes components...
   ■ Using image gcr.io/k8s-minikube/storage-provisioner:v5
   ■ Using image docker.io/kubernetesui/dashboard:v2.7.0
   ■ Using image docker.io/kubernetesui/metrics-scraper:v1.0.8
💡 Some dashboard features require the metrics-server addon. To enable all features please run:

    minikube addons enable metrics-server

🌟 Enabled addons: storage-provisioner, default-storageclass, dashboard
👉 Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
~$
```

Minikube (Cont.)

‘Minikube kubectl’ is used to run all the cluster related command. We can create the alias and rename it as the kubectl.

```
minikube kubectl -- get pods -A
```

```
alias kubectl="minikube kubectl --"
```

```
~$ alias kubectl="minikube kubectl --"
~$ kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	nginx	1/1	Running	1 (8m42s ago)	33m
kube-system	coredns-787d4945fb-9rk7g	1/1	Running	2 (8m37s ago)	160m
kube-system	etcd-minikube	1/1	Running	2 (8m42s ago)	160m
kube-system	kube-apiserver-minikube	1/1	Running	2 (8m41s ago)	160m
kube-system	kube-controller-manager-minikube	1/1	Running	2 (8m42s ago)	160m
kube-system	kube-proxy-6qv7x	1/1	Running	2 (8m42s ago)	160m
kube-system	kube-scheduler-minikube	1/1	Running	2 (8m42s ago)	160m
kube-system	storage-provisioner	1/1	Running	3 (8m42s ago)	160m
kubernetes-dashboard	dashboard-metrics-scraper-5c6664855-w99cq	1/1	Running	1 (8m42s ago)	28m
kubernetes-dashboard	kubernetes-dashboard-55c4cbbc7c-xmzvW	1/1	Running	1 (8m42s ago)	28m

```
~$
```

Namespaces in Kubernetes

- Namespaces in a Kubernetes cluster are a way to organize and isolate resources and objects within the cluster.
- A namespace provides a scope for naming and controlling the visibility of Kubernetes resources, such as Pods, Services, ConfigMaps, and Secrets.
- Each Kubernetes resource belongs to a specific namespace, and if no namespace is specified, the resource is created in the default namespace.
- By using multiple namespaces, you can logically partition the cluster into smaller, more manageable units, and avoid naming conflicts and resource collisions.


```
~$ kubectl get pods --all-namespaces
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
default	nginx	1/1	Running	1 (10m ago)	35m
kube-system	coredns-787d4945fb-9rk7g	1/1	Running	2 (10m ago)	162m
kube-system	etcd-minikube	1/1	Running	2 (10m ago)	162m
kube-system	kube-apiserver-minikube	1/1	Running	2 (10m ago)	162m
kube-system	kube-controller-manager-minikube	1/1	Running	2 (10m ago)	162m
kube-system	kube-proxy-6qv7x	1/1	Running	2 (10m ago)	162m
kube-system	kube-scheduler-minikube	1/1	Running	2 (10m ago)	162m
kube-system	storage-provisioner	1/1	Running	3 (10m ago)	162m
kubernetes-dashboard	dashboard-metrics-scraper-5c6664855-w99cq	1/1	Running	1 (10m ago)	30m
kubernetes-dashboard	kubernetes-dashboard-55c4cbbc7c-xmzvw	1/1	Running	1 (10m ago)	30m

Kubernetes Dashboard

```
--$ minikube dashboard
😄 Verifying dashboard health ...
😄 Launching proxy ...
😄 Verifying proxy health ...
🔗 Opening http://127.0.0.1:37641/api/v1/namespaces/kubernetes-dashboard/services/http:kubernetes-dashboard:/proxy/ in your default browser...
🔗 Opening in existing browser session.
```

← → 🔄 127.0.0.1:37641/api/v1/namespaces/kubernetes-dashboard/services/http:kubernetes-dashboard:/proxy/#/workloads?namespace=default

 **kubernetes** default 🔍 Search

☰ Workloads

Workloads ⓘ

- Cron Jobs
- Daemon Sets
- Deployments
- Jobs
- Pods
- Replica Sets
- Replication Controllers
- Stateful Sets

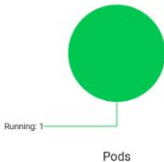
Service

- Ingresses ⓘ
- Ingress Classes
- Services ⓘ


Config and Storage

- Config Maps ⓘ
- Persistent Volume Claims ⓘ
- Secrets ⓘ
- Storage Classes

Workload Status



Pods

Name	Images	Labels	Node	Status	Rest
 nginx	nginx	run: nginx	minikube	Running	1

Few Important Minikube Commands

- minikube pause
- minikube unpause
- minikube stop
- minikube addons list
- minikube delete --all

```
$ minikube stop
```

```
👋 Stopping node "minikube" ...
```

```
🛑 1 nodes stopped.
```

```
$ minikube pause
```

```
👋 Pausing node "minikube" ...
```

```
🛑 1 nodes paused.
```

```
$ minikube delete --all
```

```
🔥 Deleting "minikube" in hyperkit ...
```

```
💔 The "minikube" cluster has been deleted.
```

```
🔥 Deleting "my-cluster" in virtualbox ...
```

```
💔 The "my-cluster" cluster has been deleted.
```


Minikube VM

```
~$ minikube ssh
```

```
$ docker images
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
nginx	latest	3f8a00f137a0	5 days ago	142MB
registry.k8s.io/kube-apiserver	v1.26.1	deb04688c4a3	3 weeks ago	134MB
registry.k8s.io/kube-scheduler	v1.26.1	655493523f60	3 weeks ago	56.3MB
registry.k8s.io/kube-controller-manager	v1.26.1	e9c08e11b07f	3 weeks ago	124MB
registry.k8s.io/kube-proxy	v1.26.1	46a6bb3c77ce	3 weeks ago	65.6MB
registry.k8s.io/etcd	3.5.6-0	fce326961ae2	2 months ago	299MB
registry.k8s.io/pause	3.9	e6f181688397	4 months ago	744kB
kubernetesui/dashboard	<none>	07655ddf2eeb	5 months ago	246MB
kicbase/echo-server	1.0	9056ab77afb8	7 months ago	4.94MB
kubernetesui/metrics-scraper	<none>	115053965e86	8 months ago	43.8MB
registry.k8s.io/coredns/coredns	v1.9.3	5185b96f0bec	8 months ago	48.8MB
registry.k8s.io/e2e-test-images/agnhost	2.39	a05bd3a9140b	8 months ago	127MB
registry.k8s.io/pause	3.6	6270bb605e12	17 months ago	683kB
gcr.io/k8s-minikube/storage-provisioner	v5	6e38f40d628d	22 months ago	31.5MB

Deployment in Kubernetes

The `kubectl create deployment` command is used to create a new Deployment in a Kubernetes cluster. A Deployment is a Kubernetes object that manages a set of identical Pods, ensuring that the desired number of replicas are running and replacing any that fail or become unresponsive.

- `kubectl create deployment hello-minikube --image=kicbase/echo-server:1.0`
- `kubectl expose deployment hello-minikube --type=NodePort --port=8080`
- `kubectl port-forward service/hello-minikube 7080:8080`
- Open in the browser: <http://localhost:7080/>

```
~$ kubectl create deployment hello-minikube --image=kicbase/echo-server:1.0
deployment.apps/hello-minikube created
~$ kubectl expose deployment hello-minikube --type=NodePort --port=8080
service/hello-minikube exposed
~$ kubectl port-forward service/hello-minikube 7080:8080
Forwarding from 127.0.0.1:7080 -> 8080
Forwarding from [::1]:7080 -> 8080
Handling connection for 7080
Handling connection for 7080
Handling connection for 7080
```

```
~$ kubectl get services hello-minikube
NAME         TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
hello-minikube NodePort    10.98.214.223 <none>        8080:32146/TCP   4m34s
~$ kubectl get deployments
NAME         READY   UP-TO-DATE   AVAILABLE   AGE
hello-minikube 1/1      1            1           5m36s
~$ kubectl get services
NAME         TYPE        CLUSTER-IP   EXTERNAL-IP   PORT(S)          AGE
hello-minikube NodePort    10.98.214.223 <none>        8080:32146/TCP   5m34s
kubernetes   ClusterIP   10.96.0.1     <none>        443/TCP          3h4m
```

Deployment in Kubernetes (Cont.)

kubectl create deployment hello-nginx --image=nginx

kubectl expose deployment hello-nginx --type=NodePort --port=80

kubectl port-forward service/hello-nginx 7000:80

Scale a resource

deployment hello-nginx will be updated to reflect the desired replicas count.

Desired replicas * 3 Actual replicas 3

This action is equivalent to: kubectl scale -n default deployment hello-nginx --replicas=3

Scale Cancel

The screenshot shows the Kubernetes Dashboard interface. On the left is a sidebar with navigation links for Workloads, Service, Config and Storage, and Cluster. The main area is titled 'Workload Status' and contains three circular progress indicators for Deployments, Pods, and Replica Sets. Below this, there are two tables: 'Deployments' and 'Pods'.

Deployments Table:

Name	Images	Labels	Pods	Created
hello-nginx	nginx	app: hello-nginx	1 / 3	a minute ago

Pods Table:

Name	Images	Labels	Node	Status	Restarts	CPU Usage (cores)	Memory Usage (bytes)	Created
hello-nginx-5cc97659f4-n7ht	nginx	app: hello-nginx pod-template-hash: 5cc97659f4	minikube	ContainerCreating	0	-	-	9 seconds ago
hello-nginx-5cc97659f4-qtpzq	nginx	app: hello-nginx pod-template-hash: 5cc97659f4	minikube	ContainerCreating	0	-	-	9 seconds ago
hello-nginx-5cc97659f4-wsqzh	nginx	app: hello-nginx pod-template-hash: 5cc97659f4	minikube	Running	0	-	-	a minute ago

To do!

- Install minikube on your computers and play
- Also go over the following interactive tutorial

<https://kubernetes.io/docs/tutorials/hello-minikube/>

YAML

XML

```
<Servers>
  <Server>
    <name>Server1</name>
    <owner>John</owner>
    <created>12232012</created>
    <status>active</status>
  </Server>
</Servers>
```

JSON

```
{
  Servers: [
    {
      name: Server1,
      owner: John,
      created: 12232012,
      status: active,
    }
  ]
}
```

YAML

```
Servers:
-   name: Server1
    owner: John
    created: 12232012
    status: active
```

YAML

- YAML, short for "YAML Ain't Markup Language", is a human-readable data format. It is often used for configuration files, data exchange between languages, and storing structured data.
- YAML is designed to be easy to read and write by humans, and is intended to be more human-friendly than other data serialization formats like JSON or XML.
- It contains key-value, lists/arrays, and dictionaries.

```
# This is a YAML document
name: John Smith
age: 30
hobbies:
  - reading
  - hiking
address:
  street: 123 Main St.
  city: Anytown
  state: CA
  zip: '12345'
```

YAML (Cont.)

List

```
# A list of fruits
fruits:
  - apple
  - banana
  - orange
```

Dictionary

```
# A dictionary of person
person:
  name: John Smith
  age: 30
  address:
    street: 123 Main St.
    city: Anytown
    state: CA
    zip: '12345'
```

Dictionary Of Lists

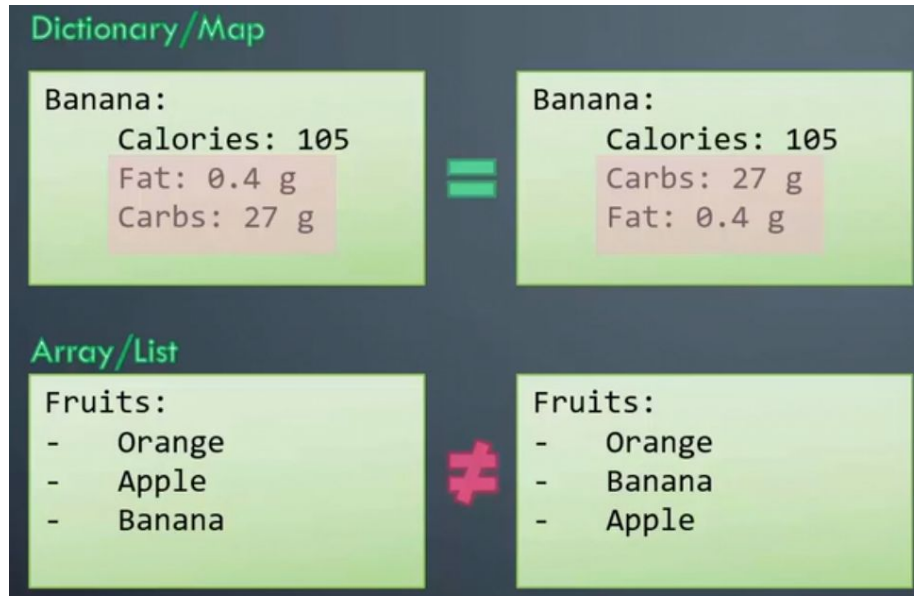
```
Fruits:
  - Banana:
      Calories: 105
      Fat: 0.4 g
      Carbs: 27 g
  - Grape:
      Calories: 62
      Fat: 0.3 g
      Carbs: 16 g
```

List of Dictionary

```
- Color: Blue
  Model:
    Name: Corvette
    Model: 1995
    Transmission : Manual
    Price: $20,000
- Color: Grey
  Model:
    Name: Corvette
    Model: 1995
    Transmission: Manual
    Price: $22,000
- Color: Red
  Model:
    Name: Corvette
    Model: 1995
    Transmission : Automatic
    Price: $20,000
- Color: Green
  Model:
    Name: Corvette
    Model: 1995
    Transmission : Manual
    Price: $23,000
```

YAML (Cont.)

In dictionary order does not matter but in list order matter.



Kubernetes YAML

In Kubernetes YAML files must contain apiVersion, kind, metadata, and spec:

pod-definition.yml

```
apiVersion: v1
kind: Pod
metadata:
  name: myapp-pod
  labels:
    app: myapp
    type: front-end
spec:
  containers:
    - name: nginx-container
      image: nginx
```

```
kubectl create -f pod-definition.yml
```

Kind	Version
POD	v1
Service	v1
ReplicaSet	apps/v1
Deployment	apps/v1

Running POD through YAML file

We can run pods through yaml files too. Here is a very simple example:

apiVersion: v1

kind: Pod

metadata:

name: redis

spec:

containers:

- name: redis

image: redis

1. Save it to file sample.yml
2. kubectl create -f example.yml
3. Kubectl get pods

Practice YAML

1. Create a pod with wrong image
2. Understand the pod error
3. Check the node it is running
4. Update the pod by fixing the error

Try the following commands:

```
kubectl set image pod redis redis=redis
```


```
kubectl edit pod redis
```

```
kubectl describe pod redis
```

```
kubectl delete pods --all
```

Replication Controllers and ReplicaSets

- Replication Controller and ReplicaSet are two different objects in Kubernetes that are used to ensure that a specified number of Pod replicas are running at all times.
- A Replication Controller is an older version of this functionality that has now been superseded by ReplicaSets
- ReplicaSet can be used for high-availability and scalability



ReplicaSet

replicaset-definition.yml

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: myapp-replicaset
  labels:
    app: myapp
    type: front-end
spec:
  template:
    metadata:
      name: myapp-pod
      labels:
        app: myapp
        type: front-end
    spec:
      containers:
        - name: nginx-container
          image: nginx
  replicas: 3
  selector:
    matchLabels:
      type: front-end
```

```
> kubectl create -f replicaset-definition.yml
```

```
replicaset "myapp-replicaset" created
```

```
> kubectl get replicaset
```

NAME	DESIRED	CURRENT	READY	AGE
myapp-replicaset	3	3	3	19s

ReplicaSet (Cont.)

```
> kubectl create -f replicaset-definition.yml
```

```
> kubectl get replicaset
```

```
> kubectl delete replicaset myapp-replicaset
```

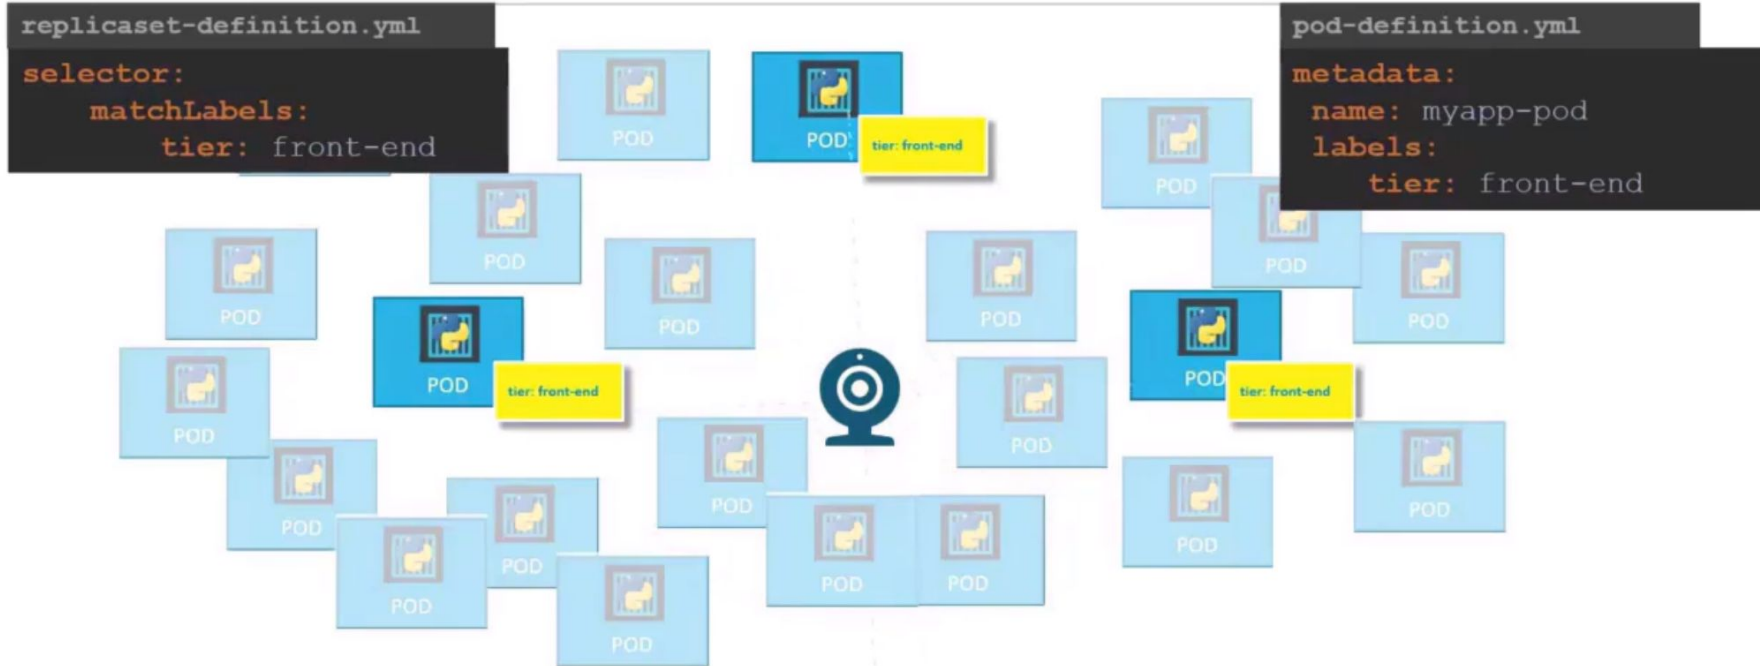
*Also deletes all underlying PODs

```
> kubectl replace -f replicaset-definition.yml
```

```
> kubectl scale --replicas=6 -f replicaset-definition.yml
```

Labels and Selectors

Labels are key-value pairs use to identify and group objects like pods, replica sets etc.



Deployments

Deployments in Kubernetes are a higher-level abstraction that allows you to manage the lifecycle of a set of Pods, using ReplicaSets under the hood.

Deployments provide several benefits over managing ReplicaSets directly:

- Rolling updates: Deployments allow you to update the image of your application gradually across the replica Pods using a rolling update strategy, which ensures that there is always a certain number of Pods running during the update process.
- Rollbacks: Deployments allow you to easily rollback to a previous version of your application if there are any issues with the new version.
- Scaling: Deployments provide an easy way to scale your application up or down by adjusting the number of replicas.
- History: Deployments keep a history of all updates made to the deployment, including the image and configuration changes, which allows you to easily track and manage the deployment's lifecycle.

Definition

```
> kubectl create -f deployment-definition.yml
```

```
deployment "myapp-deployment" created
```

```
> kubectl get deployments
```

NAME	DESIRED	CURRENT	UP-TO-DATE	AVAILABLE	AGE
myapp-deployment	3	3	3	3	21s

```
> kubectl get replicaset
```

NAME	DESIRED	CURRENT	READY	AGE
myapp-deployment-6795844b58	3	3	3	2m

```
> kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
myapp-deployment-6795844b58-5rbj1	1/1	Running	0	2m
myapp-deployment-6795844b58-h4w55	1/1	Running	0	2m
myapp-deployment-6795844b58-1fjvh	1/1	Running	0	2m

```
deployment-definition.yml
```

```
apiVersion: apps/v1
```

```
kind: Deployment
```

```
metadata:
```

```
  name: myapp-deployment
```

```
  labels:
```

```
    app: myapp
```

```
    type: front-end
```

```
spec:
```

```
  template:
```

```
    metadata:
```

```
      name: myapp-pod
```

```
      labels:
```

```
        app: myapp
```

```
        type: front-end
```

```
    spec:
```

```
      containers:
```

```
        - name: nginx-container
```

```
          image: nginx
```

```
replicas: 3
```

```
selector:
```

```
  matchLabels:
```


```
    type: front-end
```


Deployments: Update and Rollbacks


Whenever you create a new deployment or upgrade the images in an existing deployment it triggers a Rollout.

A rollout is the process of gradually deploying or upgrading your application containers. When you first create a deployment, it triggers a rollout

Revision 1



Revision 2



Deployments: Update and Rollbacks (Cont.)

```
> kubectl rollout status deployment/myapp-deployment
```

```
Waiting for rollout to finish: 0 of 10 updated replicas are available...  
Waiting for rollout to finish: 1 of 10 updated replicas are available...  
Waiting for rollout to finish: 2 of 10 updated replicas are available...  
Waiting for rollout to finish: 3 of 10 updated replicas are available...  
Waiting for rollout to finish: 4 of 10 updated replicas are available...  
Waiting for rollout to finish: 5 of 10 updated replicas are available...  
Waiting for rollout to finish: 6 of 10 updated replicas are available...  
Waiting for rollout to finish: 7 of 10 updated replicas are available...  
Waiting for rollout to finish: 8 of 10 updated replicas are available...  
Waiting for rollout to finish: 9 of 10 updated replicas are available...  
deployment "myapp-deployment" successfully rolled out
```

```
> kubectl rollout history deployment/myapp-deployment
```

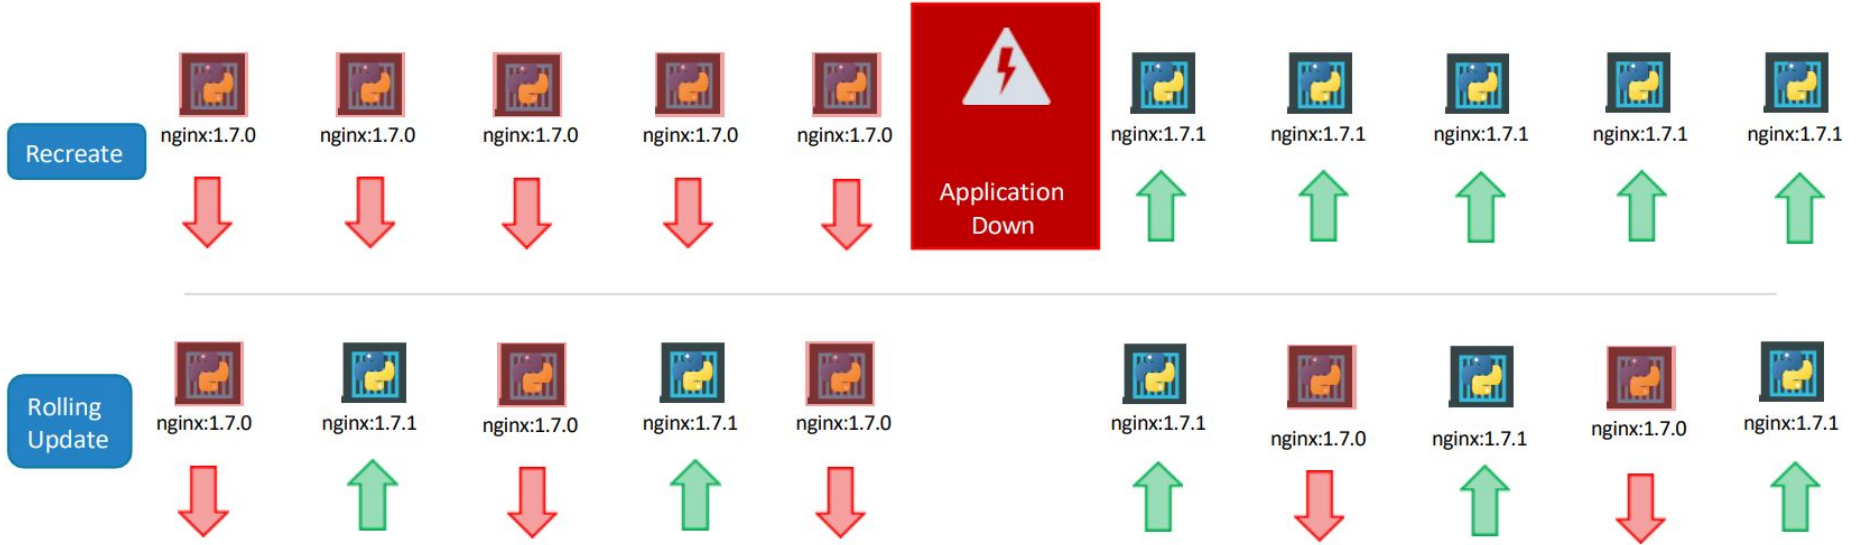
```
deployments "myapp-deployment"
```

```
REVISION  CHANGE-CAUSE
```

```
1          <none>
```

```
2          kubectl apply --filename=deployment-definition.yml --record=true
```

Recreate vs Rolling Update



Deployments (Cont.)

Create

```
> kubectl create -f deployment-definition.yml
```

Get

```
> kubectl get deployments
```

Update

```
> kubectl apply -f deployment-definition.yml
```

```
> kubectl set image deployment/myapp-deployment nginx=nginx:1.9.1
```

Status

```
> kubectl rollout status deployment/myapp-deployment
```

```
> kubectl rollout history deployment/myapp-deployment
```

Rollback

```
> kubectl rollout undo deployment/myapp-deploym
```

Services


Kubernetes Services enable communication between various components within and outside of the application. Kubernetes Services helps us connect applications together with other applications or users.



NodePort



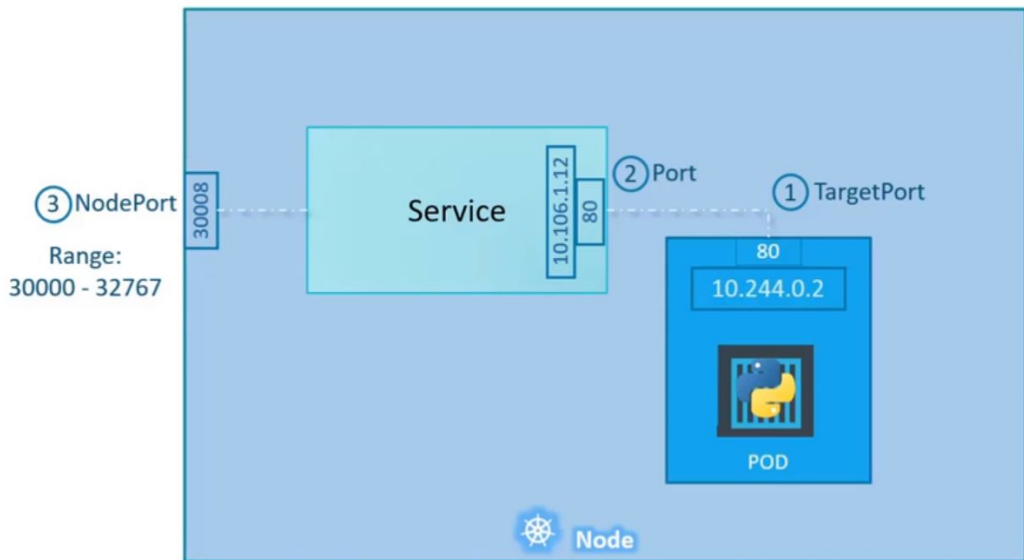
ClusterIP



LoadBalancer

Services: NodePort

NodePort is a type of service that exposes a set of pods to the external network. It allows you to access your application from outside the Kubernetes cluster by assigning a static port on each node in the cluster.



service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      *port: 80
      nodePort: 30008
```

Everything is set but target pods are missing!

Services: NodePort (Cont.)

service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: myapp-service
spec:
  type: NodePort
  ports:
    - targetPort: 80
      port: 80
      nodePort: 30008
  selector:
    app: myapp
    type: front-end
```

```
> kubectl create -f service-definition.yml
```

```
service "myapp-service" created
```

```
> kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	16d
myapp-service	NodePort	10.106.127.123	<none>	80:30008/TCP	5m

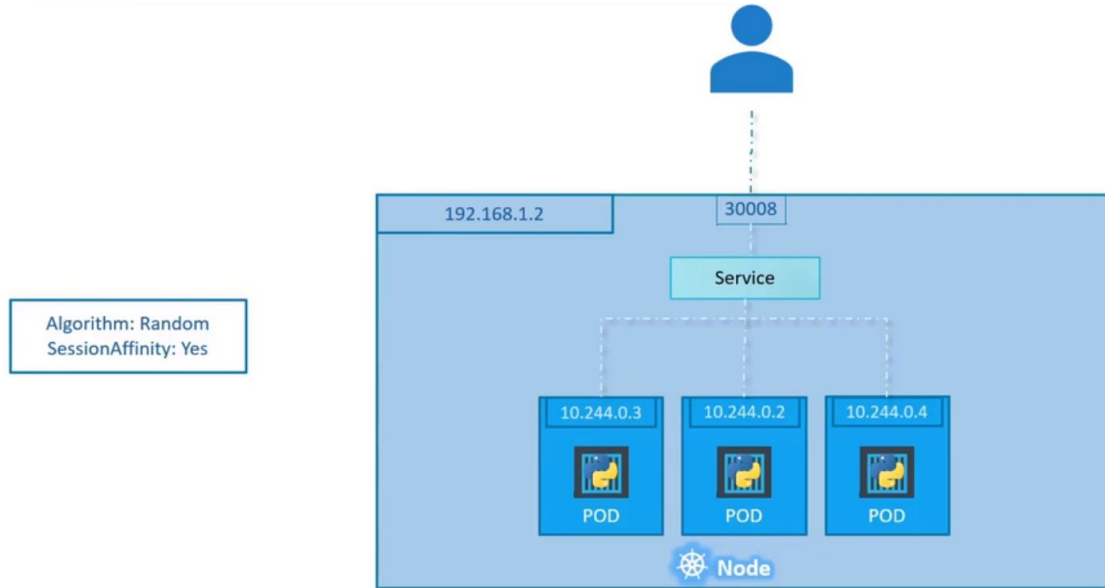
```
> curl http://192.168.1.2:30008
```

```
<html>
<head>
<title>Welcome to nginx!</title>
<style>
  body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
  }
</style>
</head>
<body>
```

Minikube service myapp-service --url gives the service URL

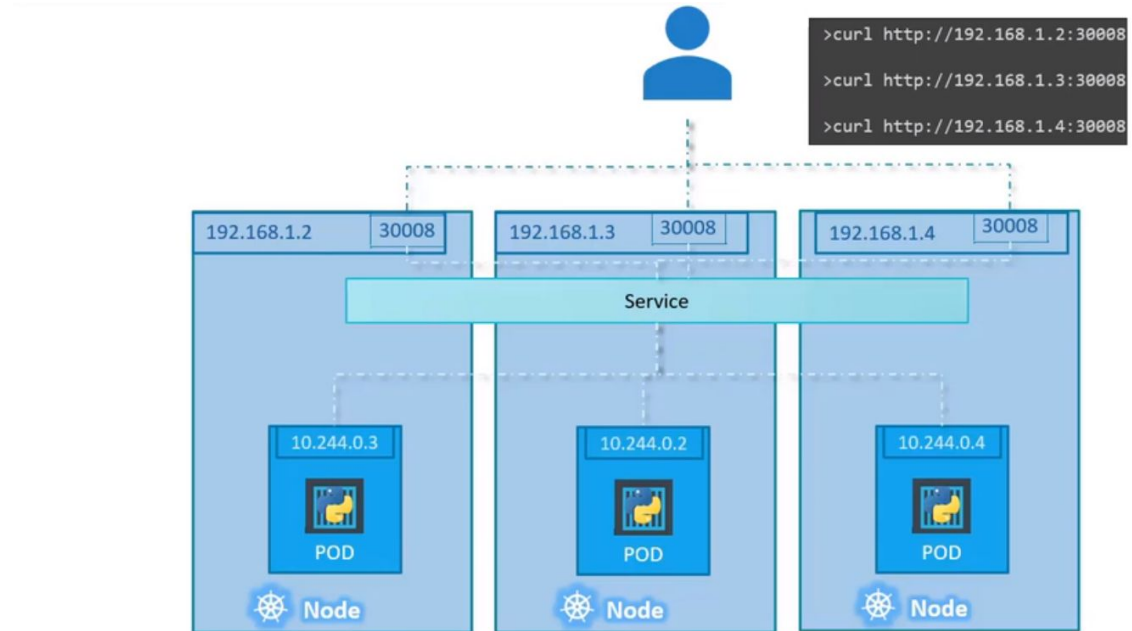
Services: NodePort (Cont.)

NodePort service automatically distribute the load to all running POD of same label.



Services: NodePort (Cont.)



On a multi node cluster, you can only need to create service and rest is handled by the Kubernetes automatically.



Services: ClusterIP

In Kubernetes, a ClusterIP Service is the default type of Service that provides a stable virtual IP address (also known as a cluster-internal IP address) that can be used to access a set of Pods running inside a cluster.

When you create a ClusterIP Service, Kubernetes assigns a virtual IP address to the Service that can be used by other Pods or Services to communicate with the Pods that are backing the Service. This virtual IP address is only accessible from within the cluster, and it is not reachable from outside the cluster.



Services: ClusterIP (Cont.)

service-definition.yml

```
apiVersion: v1
kind: Service
metadata:
  name: back-end
spec:
  type: ClusterIP
  ports:
    - targetPort: 80
      port: 80
  selector:
    app: myapp
    type: back-end
```

```
> kubectl create -f service-definition.yml
```

```
service "back-end" created
```

```
> kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	16d
back-end	ClusterIP	10.106.127.123	<none>	80/TCP	2m

Services: LoadBalancer

- In Kubernetes, a LoadBalancer Service is a type of Service that provides external access to a set of Pods running inside a cluster.
- When you create a LoadBalancer Service, Kubernetes provisions a load balancer in the cloud provider that distributes traffic to the Pods that are backing the Service.
- LoadBalancer Services are typically used when you want to expose your application to external clients, such as users accessing your application over the internet.
- By creating a LoadBalancer Service, you can provide a stable external IP address that clients can use to access your application, even if the Pods that are running the application change.

Services: LoadBalancer (Cont.)

Example voting app



<http://192.168.56.70:30035>

<http://192.168.56.71:30035>

<http://192.168.56.72:30035>

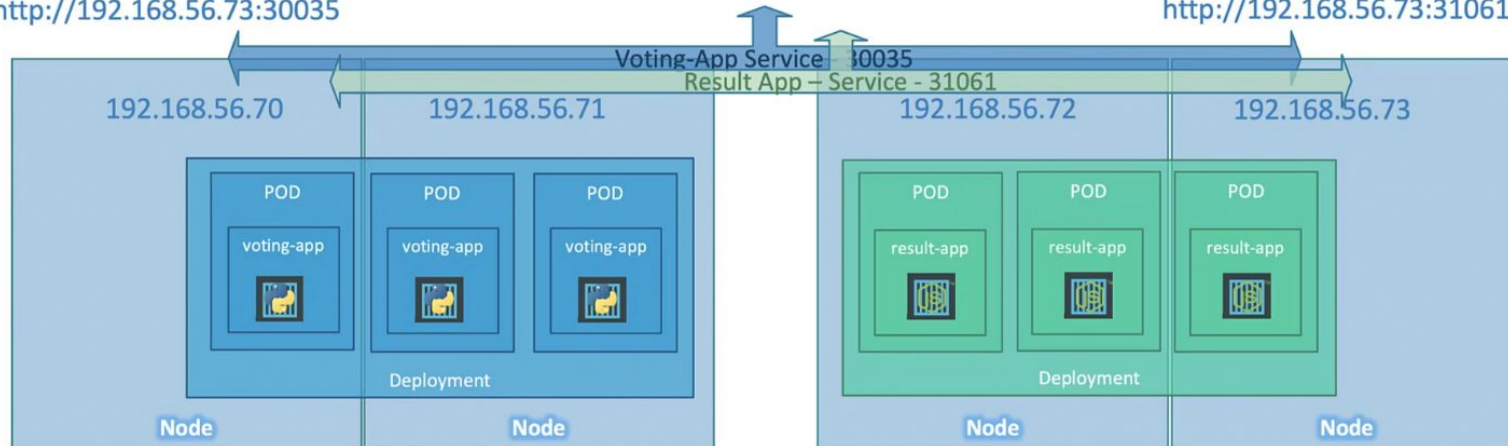
<http://192.168.56.73:30035>

<http://192.168.56.70:31061>

<http://192.168.56.71:31061>

<http://192.168.56.72:31061>

<http://192.168.56.73:31061>



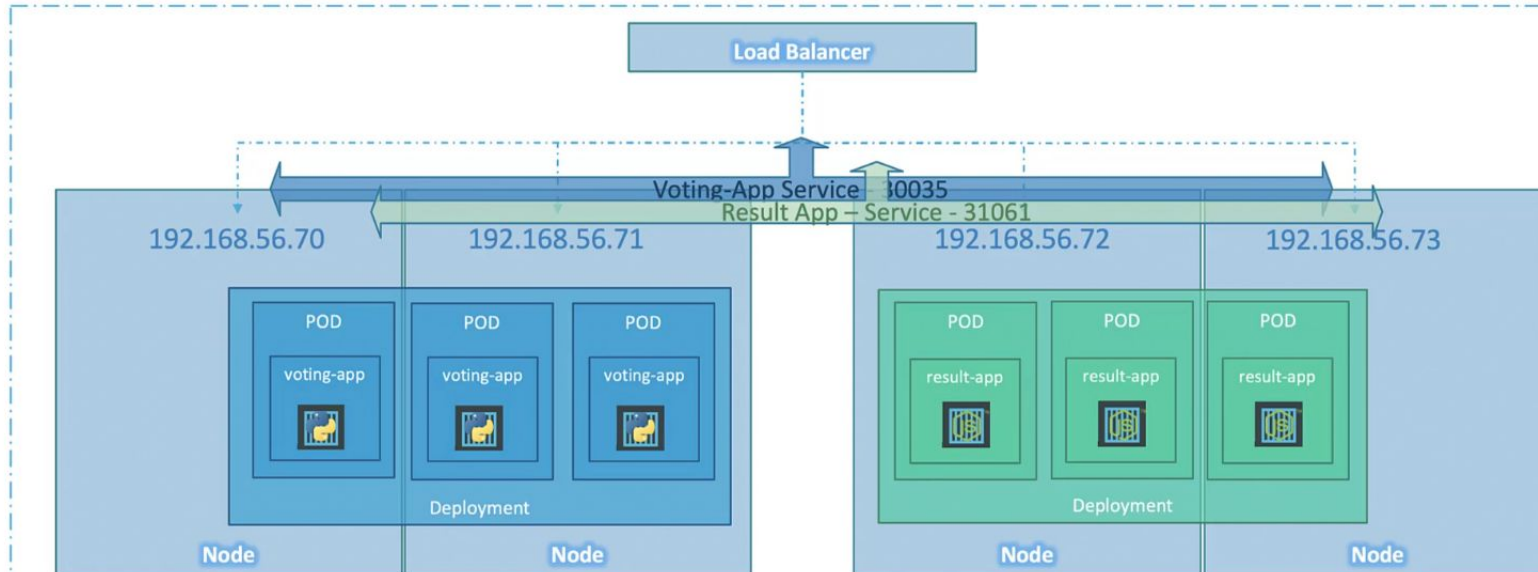
Services: LoadBalancer (Cont.)

Example voting app



<http://example-vote.com>

<http://example-result.com>



ClusterIP vs NodePort vs LoadBalancer

ClusterIP

- A ClusterIP service is an internal service that is used to provide a stable IP address and DNS name for a set of Pods.
- It is used to **enable communication** between different components of an application **running in the cluster**.
- The ClusterIP service provides a virtual IP address that can be used to access the Pods in the service, but it is **not accessible from outside the cluster**.

NodePort

- NodePort service is used to expose a service **externally by binding a port** on the node's IP address and forwarding traffic to the service.
- When you create a NodePort service, the Kubernetes API server allocates a port from a range specified by the user or defaults to a port in the range of 30000-32767. The **service is then exposed on the specified port on every node** in the cluster.

LoadBalancer

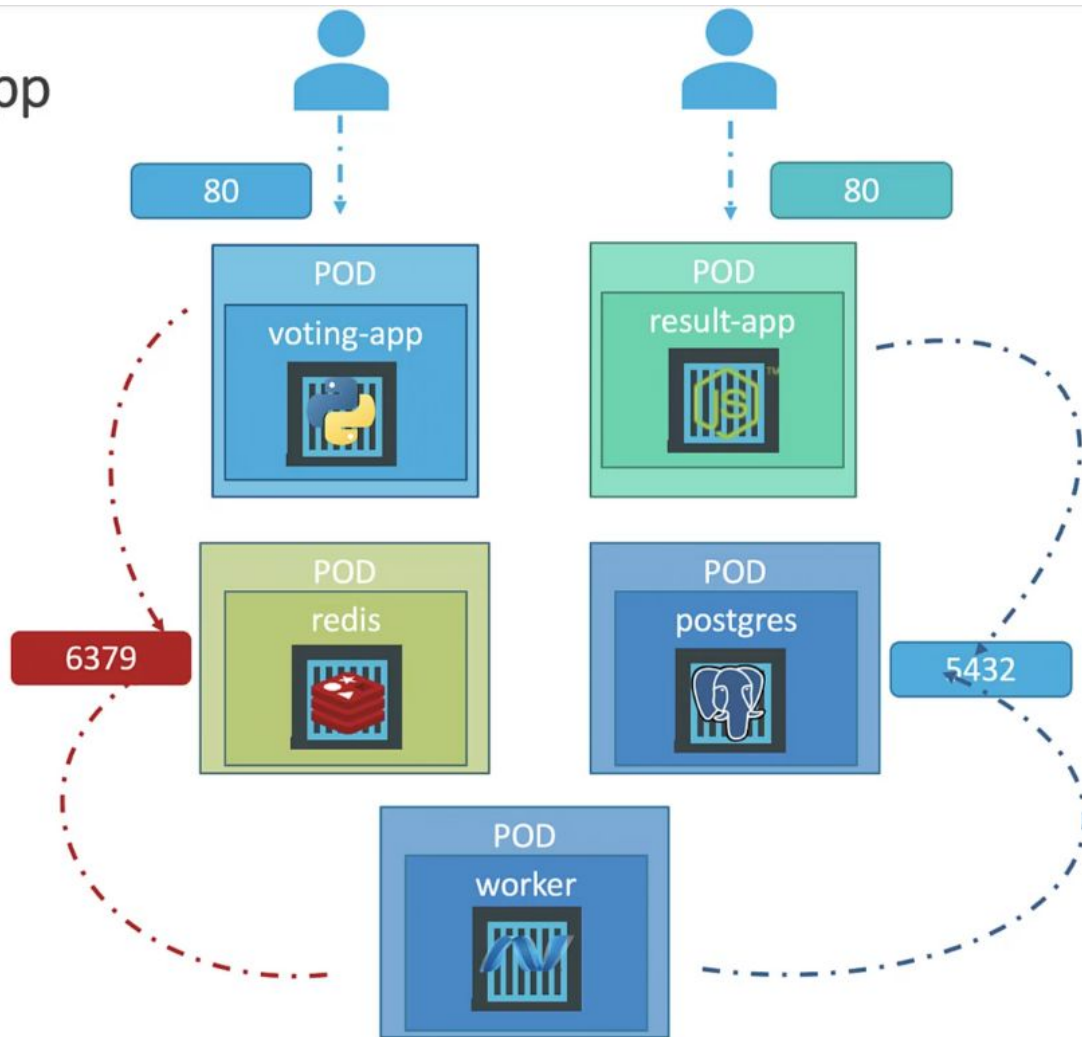
- LoadBalancer service is used to **expose a service externally** by providing a load balancer that distributes traffic to the backend Pods.
- They are typically used for production workloads that require high availability and scalability.

Example voting app

- 2. Enable Connectivity
- 3. External Access

Steps:

- 1. Deploy PODs



Example voting app

app.py

```
app = Flask(__name__)

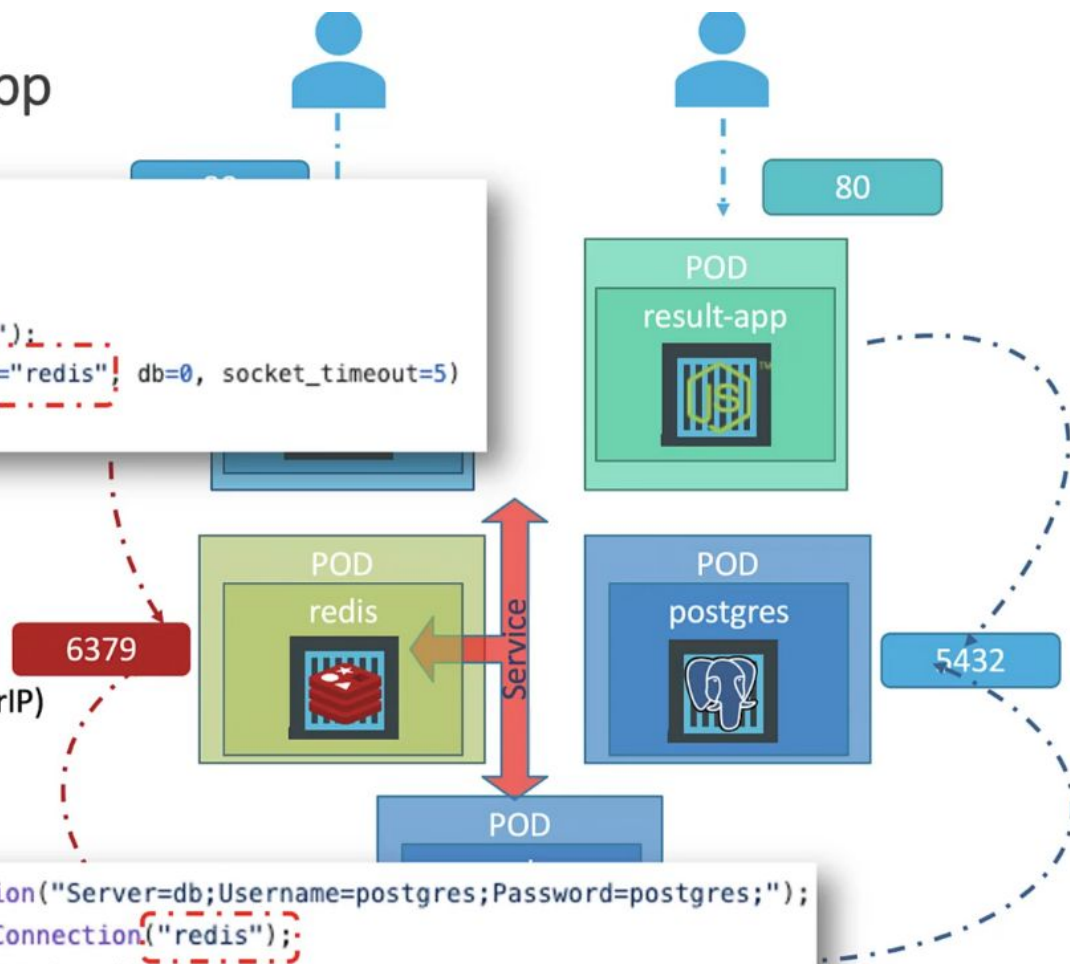
def get_redis():
    if not hasattr(g, 'redis'):
        g.redis = Redis(host="redis", db=0, socket_timeout=5)
    return g.redis
```

Steps:

1. Deploy PODs
2. Create Services (ClusterIP)
 1. redis

program.cs

```
var pgsql = OpenDbConnection("Server=db;Username=postgres;Password=postgres;");
var redisConn = OpenRedisConnection("redis");
var redis = redisConn.GetDatabase();
```



Example voting app

3. External Access

Steps:

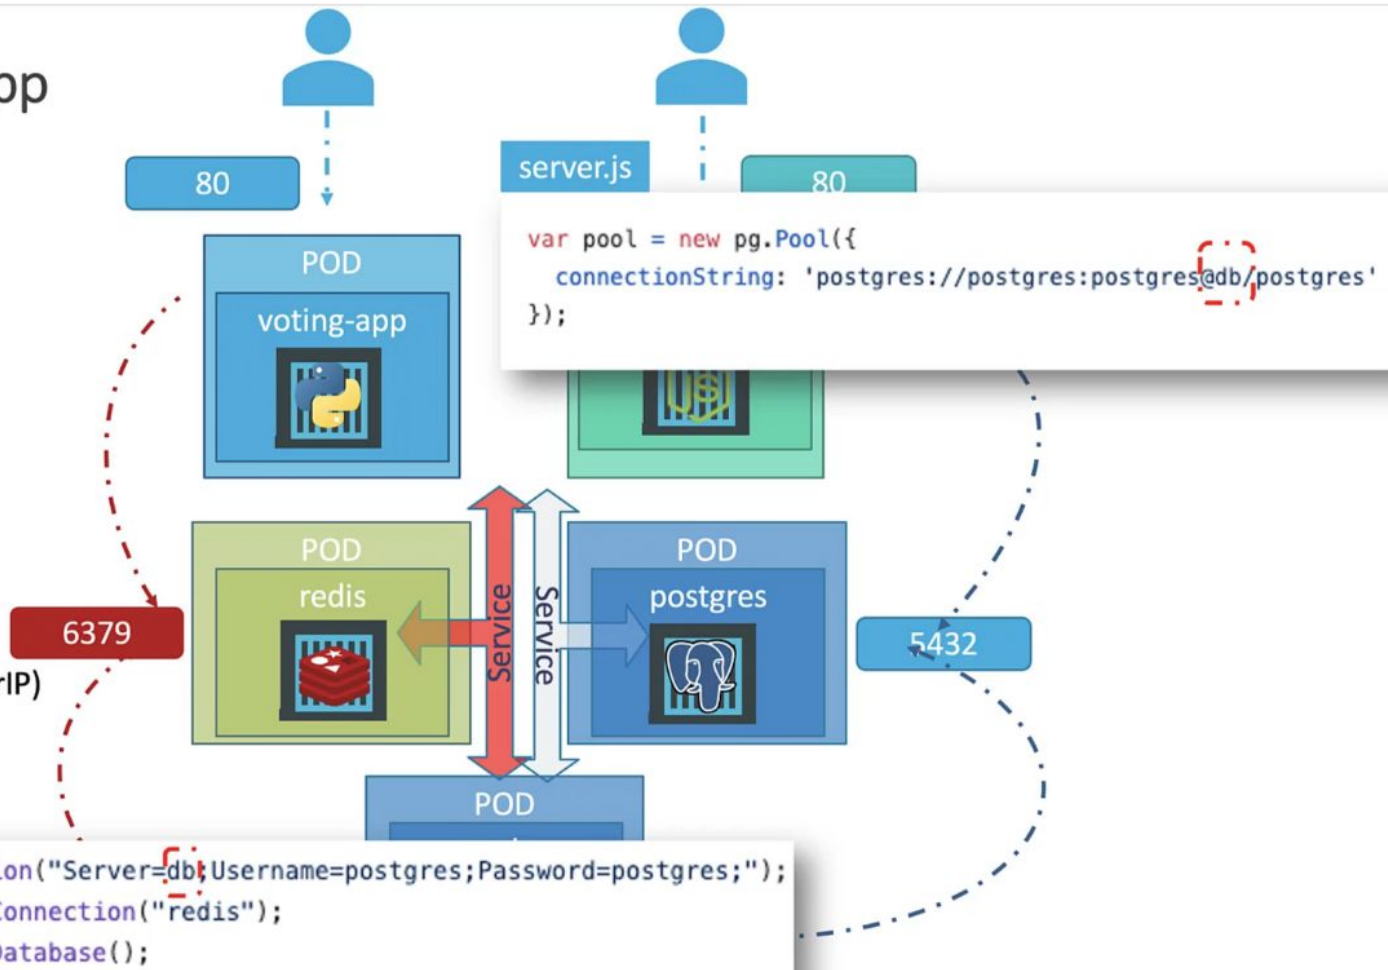
1. Deploy PODs
2. Create Services (ClusterIP)
 1. redis
 2. db

program.cs

```
var pgsq1 = OpenDbConnection("Server=db;Username=postgres;Password=postgres;");  
var redisConn = OpenRedisConnection("redis");  
var redis = redisConn.GetDatabase();
```

server.js


```
var pool = new pg.Pool({  
  connectionString: 'postgres://postgres:postgres@db/postgres'  
});
```



Example voting app

Steps:

1. Deploy PODs
2. Create Services (ClusterIP)
 1. redis
 2. db
3. Create Services (NodePort)
 1. voting-app
 2. result-app



You can add any local host image to minikube using the following command:

- `minikube cache add my-app:latest`

For any service running on minikube, you can get the URL using:

- `minikube service my-app-service --url`

If you add a local image to minikube, you might have to add the following after the image:

- `imagePullPolicy: IfNotPresent`

Credit

The material in this slide are taken from [KodeKloud.com](https://kodekloud.com)